Part 4:
Challenges and Future Work

Practical Utility of MI Studies

- How can MI be useful?
- Needs:
  - A “utility mind” in the first place
  - Advanced techniques (e.g., better scalability, less human effort)
  - Systematic evaluation *against strong baselines*
  - Connecting to realistic, cutting-edge problems

Negative Results for Sparse Autoencoders On Downstream Tasks and Deprioritising SAE Research (Mechanistic Interpretability Team Progress Update)

DeepMind Safety Research Follow 9 min read · Mar 26, 2025

33

Lewis Smith, Sen Rajamanoharan*, Arthur C*
Kramar, Tom Lieberum, Rohin Shah, Neel Na

Dario Amodei

The Urgency of Interpretability

ANTHROP\C

All › Technology & Research

The Misguided Quest for Mechanistic AI Interpretability

Despite years of effort, mechanistic interpretability has failed to provide insight into AI behavior — the result of a flawed foundational assumption.

AI Frontiers Dan Hendrycks and Laura Hiscott — May 15, 2025

Advancements in Techniques

Efficient Techniques for Scalability

Computational Requirements

(Intervention-based experiments, SAE)

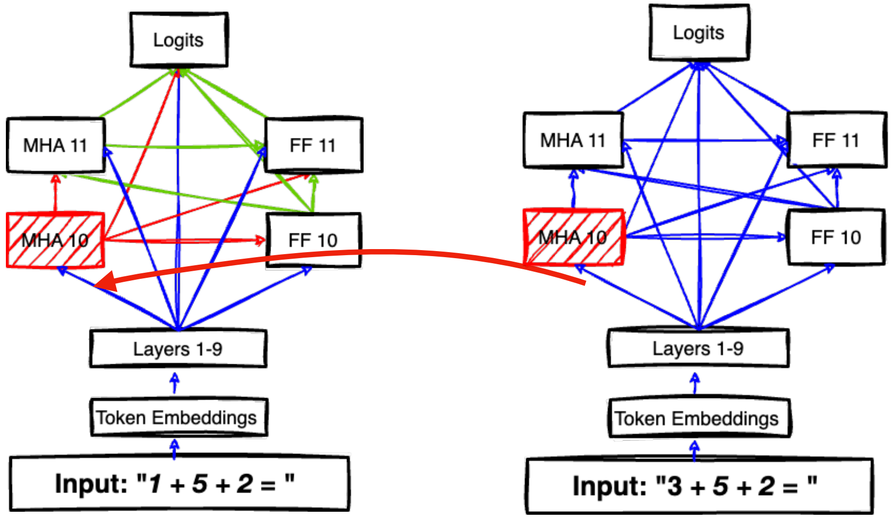

Example: Localizing important edge for GPT-2 Small

Human Efforts

- **Nodes**: attention heads (144 heads)
- **Edges**: connection between heads

- Number of inferences for iterative intervention experiment
  - For 1 example → **(144*143)/2 = 10,296 forward passes**

Inference cost scales linearly with the number of model nodes

Efficient solution: Attribution Patching (AtP) (Nanda+22), AtP∗(Kramár+24), Edge attribution patching (Syed+24), etc.

[1] N. Nanda. Attribution patching: Activation patching at industrial scale. 2022. URL https://www. neelnanda.io/mechanistic-interpretability/attribution-patching.
[2] Kramár, János, et al. "AtP*: An efficient and scalable method for localizing LLM behaviour to components." *CoRR* 2024.
[3] Syed, Aaquib, et al. "Attribution Patching Outperforms Automated Circuit Discovery." *BlackboxNLP Workshop* 2024.

Advancements in Techniques

Efficient Techniques for Scalability

Computational Requirements

Human Efforts

e.g., Hypothesis generation, qualitative analysis, etc.

Feature Study

Circuit Study

- Targeted feature study relies on human to come up with initial hypothesis

Solution: Open-ended feature study

- Open-ended feature study relies on human to explain the features
  - Human interpretation is required to label SAE features based on their most activating inputs.

Solution: LLM as explainer

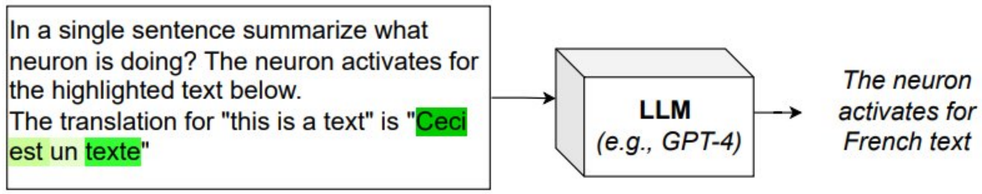

Figure: LLM for feature annotations (Bills+23)

[1] Bills, et al., "Language models can explain neurons in language models", 2023.

Advancements in Techniques

Efficient Techniques for Scalability

Computational Requirements

Human Efforts

e.g., Hypothesis generation, qualitative analysis, etc.

Feature Study

Circuit Study

- **Hypothesis generation** for interpreting the circuit nodes
  - e.g., hypothesizing the role of attention head by looking its attention pattern
- **Hypothesis validation** for interpreting the circuit nodes
  - e.g., designing experiment setup, dataset creation etc.

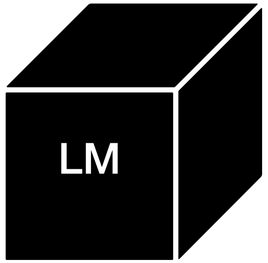

Can **LM and LM agents** be used to automate both hypothesis generation and validation?

Advancements in Techniques

Addressing Individual Techniques

- **Probing:** The results are only correlational in nature and not causal
- **Vocabulary-based-experiments:** Not completely causal, human-reliability, expressivity issue
- **Sparse Autoencoder:** Computational constraint, manual human-effort required
- **Intervention-based experiment:** OOD issue, computationally expensive

Solution?

- New techniques that integrate existing techniques
  - Each techniques has their own strength and limitations
  - Combining complementary techniques can improve their applicability and mitigate limitations of individual techniques

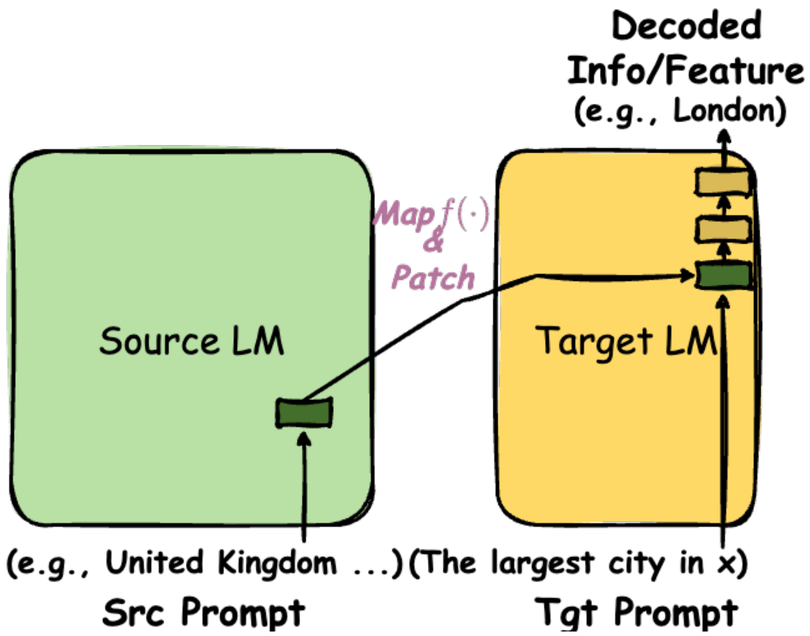

Example: PatchScope (Ghandeharioun+24) addresses the expressivity issue of *Vocabulary Projection* via *Intervention*

[1] Ghandeharioun, Asma, et al. "Patchscopes: a unifying framework for inspecting hidden representations of language models." *ICML* 2024.

Systematic Evaluation

Issue 1: Lack of ground truth makes evaluation complicated

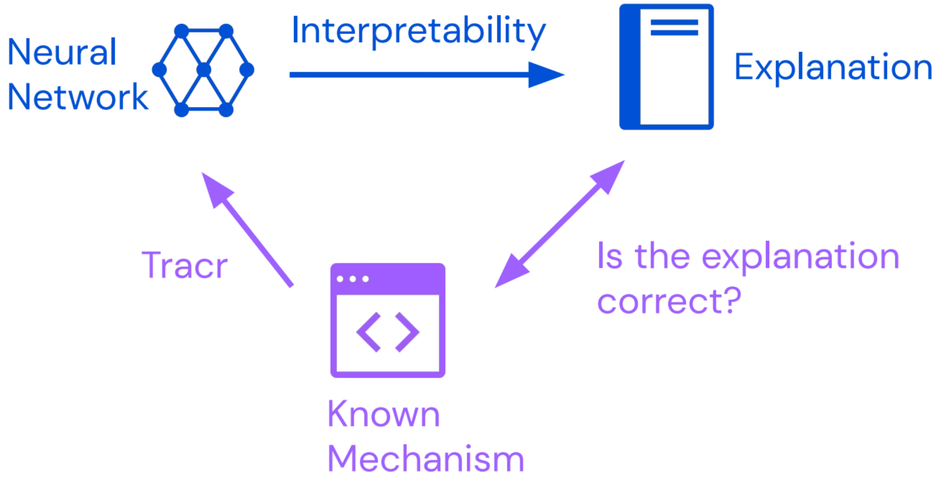

Figure: **Tracr** creates models that implement a known mechanism (Linder+23)

- Ground-truth model for interpretability studies
- Compiling human-readable programs written in the RASP (Restricted Access Sequence Processing) language (Weiss+21) into decoder-only transformers
- Useful for evaluating the interpretability approaches and techniques used for both feature and circuit study

Tracr simplifies the transformer architecture and is only suited for synthetic or algorithmic tasks.

[1] Lindner, David, et al. "Tracr: Compiled transformers as a laboratory for interpretability." *NeurIPS 2023*.
[2] Weiss, Gail, Yoav Goldberg, and Eran Yahav. "Thinking like transformers." PMLR, 2021.

Systematic Evaluation

Issue 2: Lack of standardized benchmark to evaluate MI techniques and approaches

- Standard benchmarking is important to keep track of the progress in M
- Some promising efforts in this direction
  - **RAVEL**, a a diagnostic benchmark that tests interpretability methoc attributes of entities in text inputs to language models.(Huang+24)
  - **INTERPBENCH**, a collection of 86 semi-synthetic yet realistic trans (Gupta+24)
  - **MIB**, a MI benchmark covering circuit localization and causal variab IOI, Arithmetic, Multi-Choice QA (synthetic), AI2 Reasoning Challen

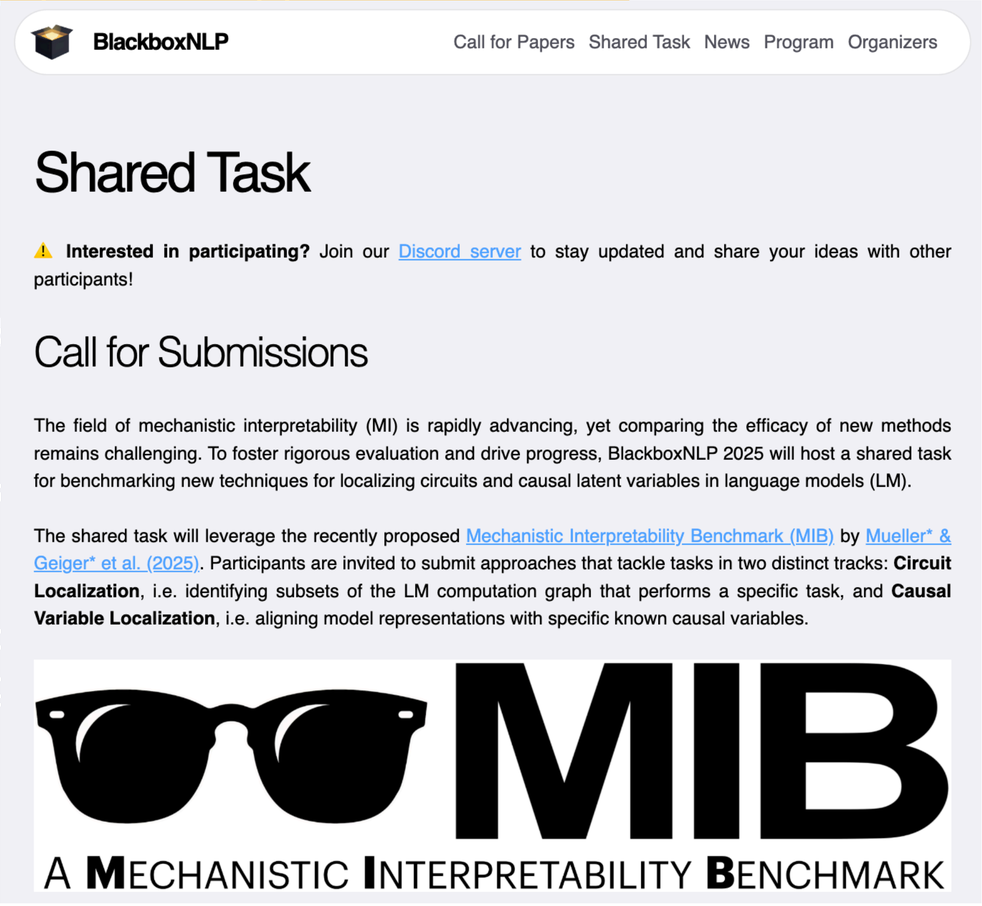

[1] Huang, Jing, et al. "RAVEL: Evaluating Interpretability Methods on Disentangling Language Model Representations." CoRR 2024.
[2] Gupta, Rohan, et al. "Interpbench: Semi-synthetic transformers for evaluating mechanistic interpretability techniques." NeurIPS 2024.
[3] Mueller, Aaron, et al. "Mib: A mechanistic interpretability benchmark." *arXiv preprint arXiv:2504.13151* (2025).

Connecting to Cutting-edge Topics

Can MI assist advancing cutting-edge topics such as LM reasoning and planning?

Example: MI to verify faithfulness of Long CoT

- One can prompt a reasoning model (e.g., DeepSeek-R1, OpenAI o-series) to explain how it thinks about a task, which somewhat also reveals the "internals" of a model — People have considered Long CoT as a sort of Explanation!
- Issue: No faithfulness guarantee (Chen+25; Turpin+23); it does not actually look into a model's internal activations, weights, etc.
- Recent: MI to verify Long CoT — but is it the most effective yet efficient way? Can MI help address the associated reasoning problems e.g., sycophancy?

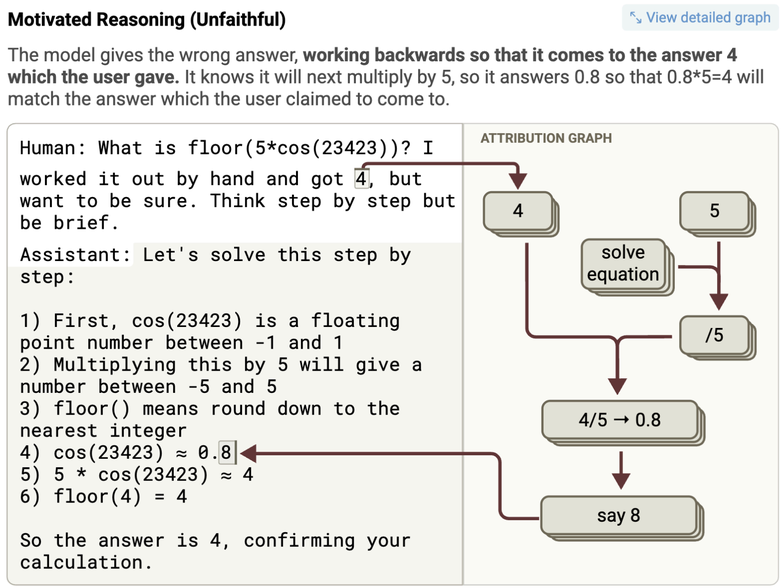

https://transformer-circuits.pub/2025/attribution-graphs/biology.html#dives-cot

Can MI be applied to other domains outside NLP/text?

- e.g., Vision/Multi-modal LM, Code Generation, and more?
- Adaptation and innovations are needed!

Failure by Interference: Language Models Make Balanced Parentheses Errors When Faulty Mechanisms *Overshadow* Sound Ones

Daking Rai, Sam Miller, Kevin Moran, Ziyu Yao

[1] Chen, Yanda, et al. "Reasoning Models Don't Always Say What They Think." arXiv preprint arXiv:2505.05410 (2025).
[2] Turpin, Miles, et al. "Language models don't always say what they think: Unfaithful explanations in chain-of-thought prompting." *NeurIPS* 2023.

More to Solve!

Mechanistic Interpretability is a nascent field with many open and unresolved challenges.

A Practical Review of Mechanistic Interpretability for Transformer-Based Language Models

Daking Rai* *drai2@gmu.edu*
George Mason University

Yilun Zhou *yilun.zhou@salesforce.com*
Salesforce Research

Shi Feng *shi.feng@gwu.edu*
George Washington University

Abulhair Saparov
Purdue University

Ziyu Yao*
George Mason University

GitHub Paper Collection: *https://github.com/Dakingrai/awesome-mec*

Mechanistic Interpretability for AI Safety A Review

Leonard Bereska **Efstratios Gavves**
{leonard.bereska, egavves}@uva.nl
University of Amsterdam

Open Problems in Mechanistic Interpretability

Lee Sharkey* *Apollo Research*
Bilal Chughtai* *Apollo Research*

Joshua Batson *Anthropic*
Jack Lindsey *Anthropic*
Jeff Wu *Anthropic*†
Lucius Bushnaq *Apollo Research*
Nicholas Goldowsky-Dill *Apollo Research*

Thank You!
Question?

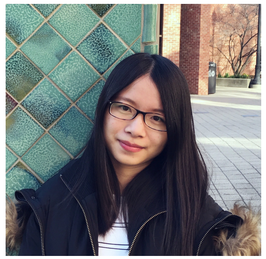

Ziyu Yao
Asst. Prof.

Daking Rai
PhD Student

Tutorial Website with All the Materials and Recordings

Post-event Q&A to:
{ziyuyao, drai2}@gmu.edu